



Exercise 11.1: Service Mesh

If you have a large number of services to expose outside of the cluster, or to expose a low-number port on the host node you can deploy an ingress controller. While nginx and GCE have controllers mentioned a lot in Kubernetes.io, there are many to choose from. Even more functionality and metrics come from the use of a service mesh, such as Istio, Linkerd, Contour, Aspen, or several others.

1. We will install linkerd using their own scripts. There is quite a bit of output. Instead of showing all of it the output has been omitted. Look through the output and ensure that everything gets a green check mark. Some steps may take a few minutes to complete. Each command is listed here to make install easier. As well these steps are in the `setupLinkerd.txt` file.

```
student@cp:~$ curl -sL run.linkerd.io/install | sh

student@cp:~$ export PATH=$PATH:/home/student/.linkerd2/bin

student@cp:~$ echo "export PATH=$PATH:/home/student/.linkerd2/bin" >> $HOME/.bashrc

student@cp:~$ linkerd check --pre

student@cp:~$ linkerd install | kubectl apply -f -

student@cp:~$ linkerd check

student@cp:~$ linkerd viz install | kubectl apply -f -

student@cp:~$ linkerd viz check

student@cp:~$ linkerd viz dashboard &
```

2. By default the GUI is on available on the localhost. We will need to edit the service and the deployment to allow outside access, in case you are using a cloud provider for the nodes. Edit to remove all characters after equal sign for `-enforced-host`, which is around line 59.

```
student@cp:~/app2$ kubectl -n linkerd-viz edit deploy web
```

YAML

```
spec:
  containers:
  - args:
    - -linkerd-controller-api-addr=linkerd-controller-api.linkerd.svc.cluster.local:8085
    - -linkerd-metrics-api-addr=metrics-api.linkerd-viz.svc.cluster.local:8085
    - -cluster-domain=cluster.local
    - -grafana-addr=grafana.linkerd-viz.svc.cluster.local:3000
    - -controller-namespace=linkerd
    - -viz-namespace=linkerd-viz
    - -log-level=info
    - -enforced-host=                                #<-- Remove everything after equal sign
    image: cr.l5d.io/linkerd/web:stable-2.10.2
    imagePullPolicy: IfNotPresent
```

3. Now edit the http nodePort and type to be a NodePort.

```
student@cp:~/app2$ kubectl edit svc web -n linkerd-viz
```

YAML

```

1 ....
2 ports:
3   - name: http
4     nodePort: 31500                                #<-- Add line with an easy to remember port
5     port: 8084
6   ....
7   sessionAffinity: None
8   type: NodePort                                    #<-- Edit type to be NodePort
9   status:
10    loadBalancer: {}
11   ....

```

4. Test access using a local browser to your public IP. Your IP will be different than the one shown below.

```
student@cp:~/app2$ curl ifconfig.io
```

```
1 104.197.159.20
```

5. From you local system open a browser and go to the public IP and the high-number nodePort.

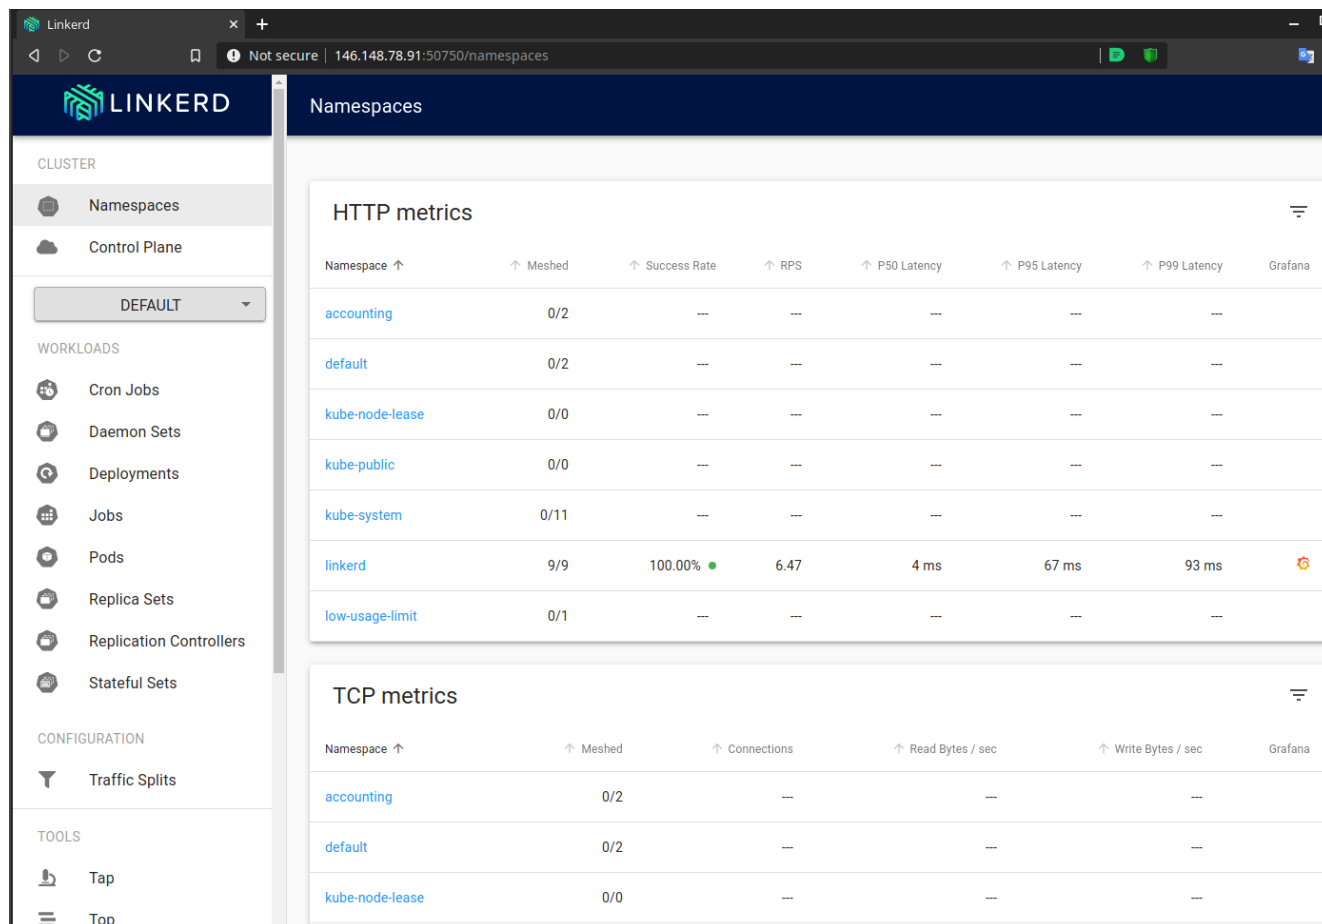


Figure 11.3: Main Linkerd Page

6. In order for linkerd to pay attention to an object we need to add an annotation. The **linkerd inject** command will do this for us. Generate YAML and pipe it to **linkerd** then pipe again to **kubectl**. Expect an error about how the object was created, but the process will work. The command can run on one line if you omit the back-slash.

```
student@cp:~/app2$ kubectl -n accounting get deploy nginx-one -o yaml | \
linkerd inject - | kubectl apply -f -
```

```
1 <output_omitted>
```

7. Check the GUI, you should see that the accounting namespaces and pods are now meshed, and the name is a link.

8. Generate some traffic to the pods, and watch the traffic via the GUI. Use the service-lab service.

```
student@cp:~$ kubectl -n accounting get svc
```

```
1 NAME          TYPE        CLUSTER-IP    EXTERNAL-IP    PORT(S)        AGE
2 nginx-one     ClusterIP   10.107.141.227 <none>         8080/TCP        5h15m
3 service-lab   NodePort    10.102.8.205   <none>         80:30759/TCP    5h14m
4
```

```
student@cp:~$ curl 10.102.8.205
```

```
1 <!DOCTYPE html>
2 <html>
3 <head>
4 <title>Welcome to nginx!</title>
5 <output_omitted>
```

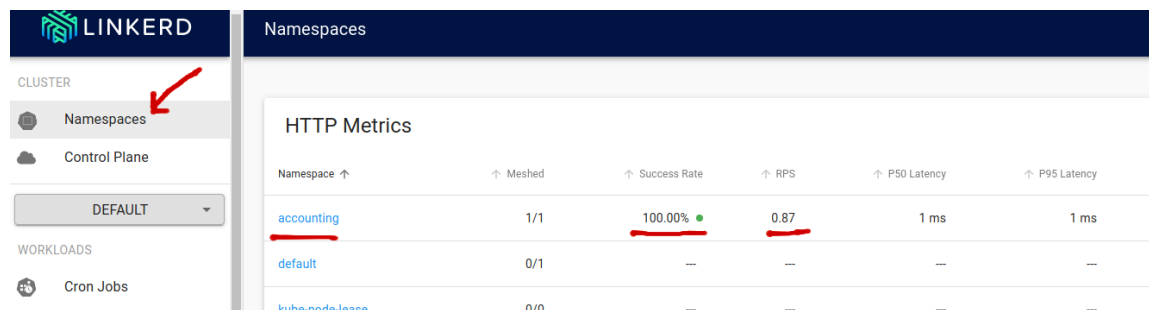


Figure 11.4: Now shows meshed

9. Scale up the nginx-one deployment. Generate traffic to get metrics for all the pods.

```
student@cp:~$ kubectl -n accounting scale deploy nginx-one --replicas=5
```

```
1 deployment.apps/nginx-one scaled
```

```
student@cp:~$ curl 10.102.8.205 #Several times
```

10. Explore some of the other information provided by the GUI. Note that the initial view is of the default namespaces. Change to accounting to see details of the nginx-one deployment.

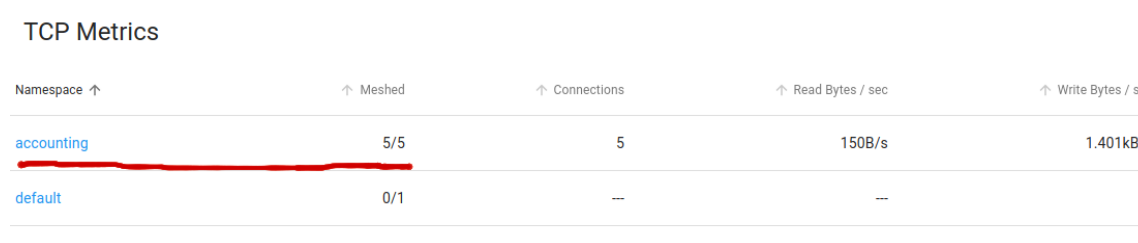


Figure 11.5: Five meshed pods